

Elastic ELK

Elasticsearch, Logstash & Kibana

Search & Analyze log files in clusters & Real Time

Berney Léonard, Bron Sacha, Minder Valentin, Salathe Fabien
HEIG-VD - CLD - 2016



elastic



logstash



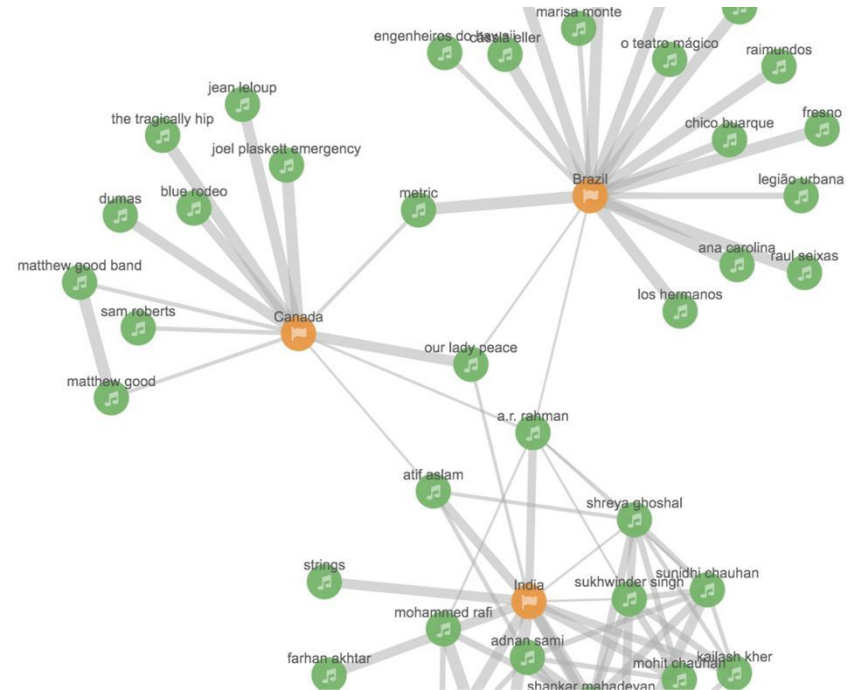
Table of contents

- Introduction
- Stack ELK = Elastic + Logstash + Kibana
- Logstash
- Elastic
 - Demo 1: How to use Elastic Search
- Kibana
 - Demo 2: Visualize Swiss Public Transportation Stops
- Elasticsearch on AWS
- Conclusion

Introduction

Elasticsearch is a **search server**.

Elasticsearch is **distributed** and frequently used via a **web page**.



Stack "ELK"

- Your existing servers
- Logstash
- Elasticsearch
- Kibana

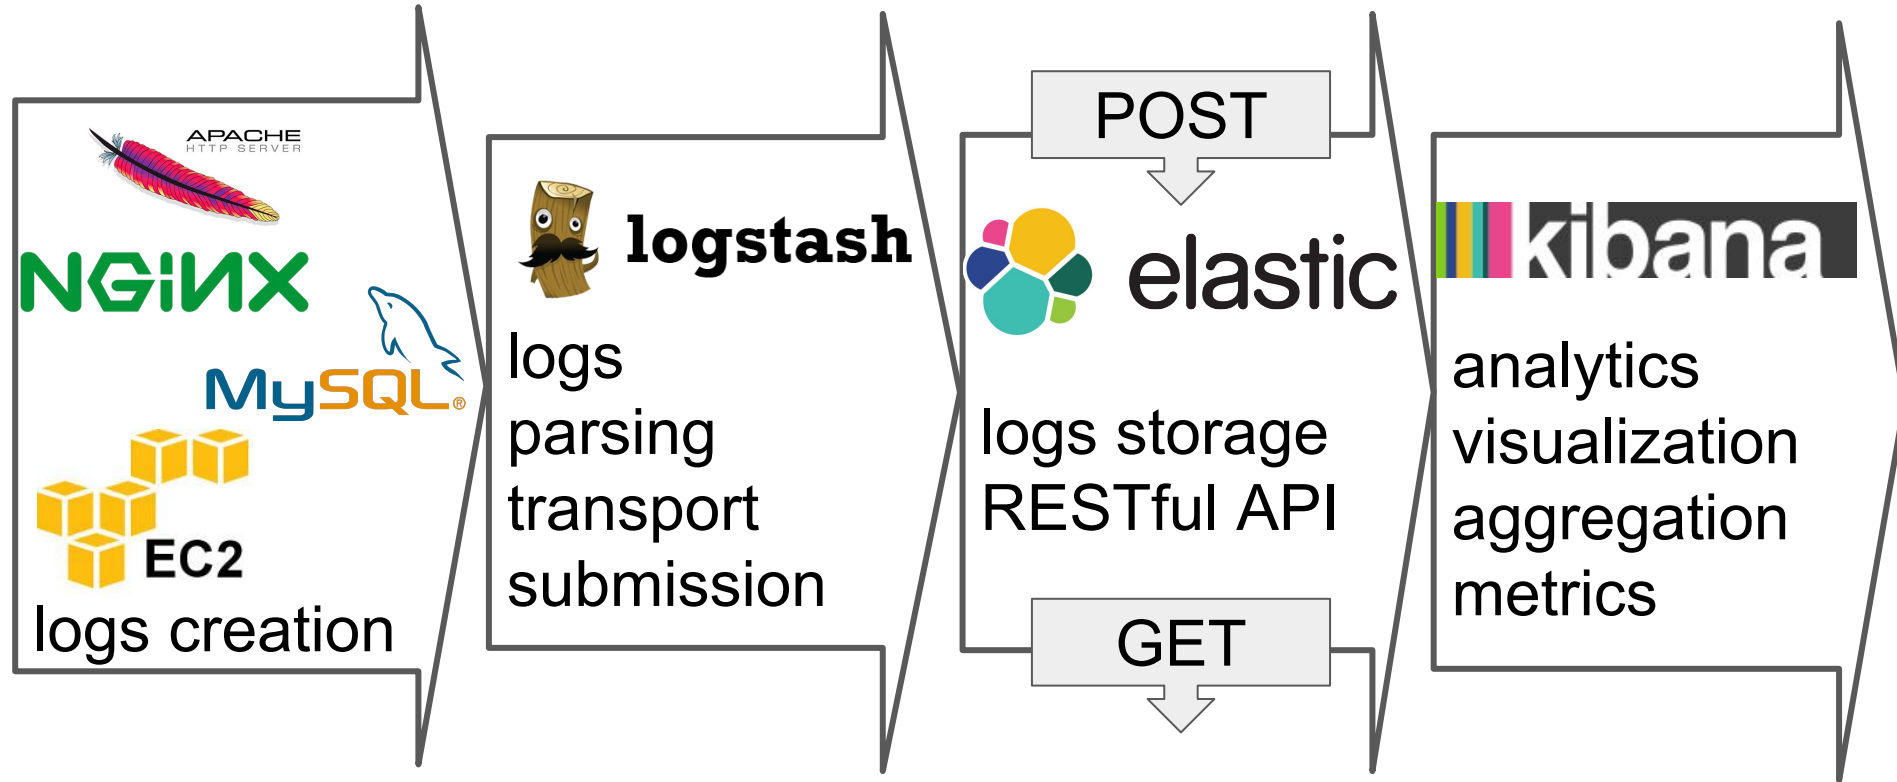


Stack "ELK"

- Logstash: abstraction and wrappers with for **logs parsing**
- Elasticsearch: **RESTful API** for log storage & retrieval
- Kibana: real-time data gathering, **visualization** and aggregation



Stack "ELK" in context





logstash

Logstash is a tool used to collect, handle and the push events and log messages to the server.

- Logs input can be customised to "watch" any log file.
- Filters can be created to handle any type of log structure.
- Processed information are then push to the server, Elasticsearch.



logstash

logstash.conf

```
input {
  file {
    path => "/var/log/apache2/access_log"
    type => "apache_access_log"
  }
  file {...}
}

filter {
  if [type] == "apache_access_log" {
    mutate { replace => { "type" => "apache-access" } }
    grok {
      match => { "message" => "%{COMBINEDAPACHELOG}" }
    }
    date {
      match => [ "timestamp", "dd/MMM/yyyy:HH:mm:ss Z" ]
    }
  }
  if [type] == "apache_error_log" {...}
  if [type] == "syslog" {...}
}

output {
  elasticsearch { hosts => "search-elastic-search-heig-superman.amazonaws.com:80" }
}
```

```
patterns.d/apache-error
APACHE_ERROR_LOG \[(?<timestamp>%{DAY:day} %{MONTH:month} %{MONTHDAY}
%{TIME} %{YEAR})\] \[.*%{LOGLEVEL:loglevel}\] \[pid %{NUMBER:pid}\] \
[client %{IP:clientip}.*\] %{GREEDYDATA:message}
```




How it works ?

- Ready and easy-to-use and simple RESTful JSON API
- Stores everything
- Retrieves everything, selected fields, or aggregations (sum, max, avg. etc)
- Accepts any input within the PUT (fields don't need to be declared)
- PUT on `server.com/index/ressource/` creates a NEW ressource under index pattern (a single server may have several indexes)
- By default, no security, no authentication (plugin/add-on available)
- OK for log display, analysis & visualisation, not sensitive data.



Advantages

- RESTful JSON API
- Fast deployment
- Aggregation
- Designed to be distributed
- Open source
- Query DSL syntax is really flexible and easy to use: general search or specific field / choose what fields to output / usual condition (AND, OR, NOT, ranges)
- even easier with kibana visualisation
- Official Javascript npm package



Drawbacks

- Still relatively **new project**, not a lot of documentations, maybe not very stable.
- **Security**: Elasticsearch does not provide any authentication or access control functionality.
- **Transactions**: There is no support for transactions or processing on data manipulation
- **Durability**: ES is distributed and fairly stable but backups and durability are not as high priority as in other data stores. This is probably the most important if you're going to make ES the primary store since losing your data is never good.



```
curl -XPUT 'https://example.com/logs/windows/bug12421?pretty' -d '{ "err" :  
"Path too long" }'
```

```
curl -XGET 'https://example.com/logs/windows/_search?q=Path&pretty=true'  
{  
  "took" : 3, "timed_out" : false, "_shards" : { "total" : 5, "successful" : 5, "failed" : 0 },  
  "hits" : {  
    "total" : 1, "max_score" : 0.15342641, "hits" : [{  
      "_index" : "logs",  
      "_type" : "windows",  
      "_id" : "bug12421",  
      "_score" : 0.15342641,  
      "_source": { "err" : "Path too long" }  
    }]  
  }  
}
```



Demo 1: How to Use Elasticsearch

- Servers (Apache, database, system ...) send their data to the ELS
- Create some errors/logs (db shutdown, wrong login, kill processes etc)
- Explore the logs to find the issues
- The same ELS contains logs from multiple sources
- The logs update automatically & regularly

https://search-elastic-search-heig-3nhbodzwhflo56pew23jotan6a.eu-central-1.es.amazonaws.com/_plugin/kibana/#



Aggregation & Visualisation Software

- Discover : search & aggregate
 - search anywhere
 - search by field, toggle arguments
 - view the field you choose
- Visualisation : charts
 - pie chart, area chart, line chart, vertical chart, metrics, tile map etc.
 - from any (cross-) search possible
- Dashboard : home screen for saved visualisations



Los Angeles Police Department Crime Dashboard (source: www.elastic.co)



Demo 2: Aggregation & Visualisation of Public Transportation Stop

- raw set from Opendata
- little ameliorations (formats of dates, numbers, GPS coordinates)
- Interesting fields
 - Stop name (Yverdon*): `Nom:Yverdon*`
 - Municipality (Commune): `NomCommune:Lausanne`
 - Company (CFF, CarPostal, TL, TPG...): `AbreviationET:sbbcfffs`
 - Type (bus, cheminfer, tram, telepherique, bateau, funiculaire, cremailere, metro): `MoyenTransport`
 - Type d'exploitation (Arret/PointChargement): `TypePointExploitation`
 - Altitude above sea level: `Altitude:[1000 TO *], Altitude:>=500`
 - GPS Coordinate: `location`

xtf_id	ch14uvag00064399	#unique ID
Numero	8504200	#Stop
Nom	Yverdon-les-Bains	Full stop name
Abreviation	YV	Abbreviation
RespDonneesAbreviation	SBBCFFFFS	Entreprise: + empty
NumeroET	1	#Entreprise: [1-9999]
AbreviationET	SBBCFFFFS	Entreprise: pag (CarPostal), sbbcffffs, vbz, tpg, tpf, tl, vzo, trn, bvb, zvb, bos, vbg, svb, blt, rbs, bls, sti, fr, rhb, vr, etc...
TypePointExploitation	Arret_et_PointChargement	Type: Arret / PointChargement / PointExploitationSimple
MoyenTransport	CheminFer	Type: bus, cheminfer, tram, telepherique, bateau, funiculaire, cheminfercremailiere, metro
Altitude	435	Meters ASL: [186-3454]
NumeroCommune	5938	OFS Number [1-9999]
NomCommune	Yverdon_les_Bains	Municipality
y_Coord_Est, x_Coord_Nord	539100, 181500	CH1903/LV03 - Swiss Army Coordinates
location	{"lat":46.7818,"lon":6.6411}	WGS84 GPS Coordinate (computed from CH1903*)

Other date info, not shown: DebutValidite, FinValidite, DateTraitement, Etat

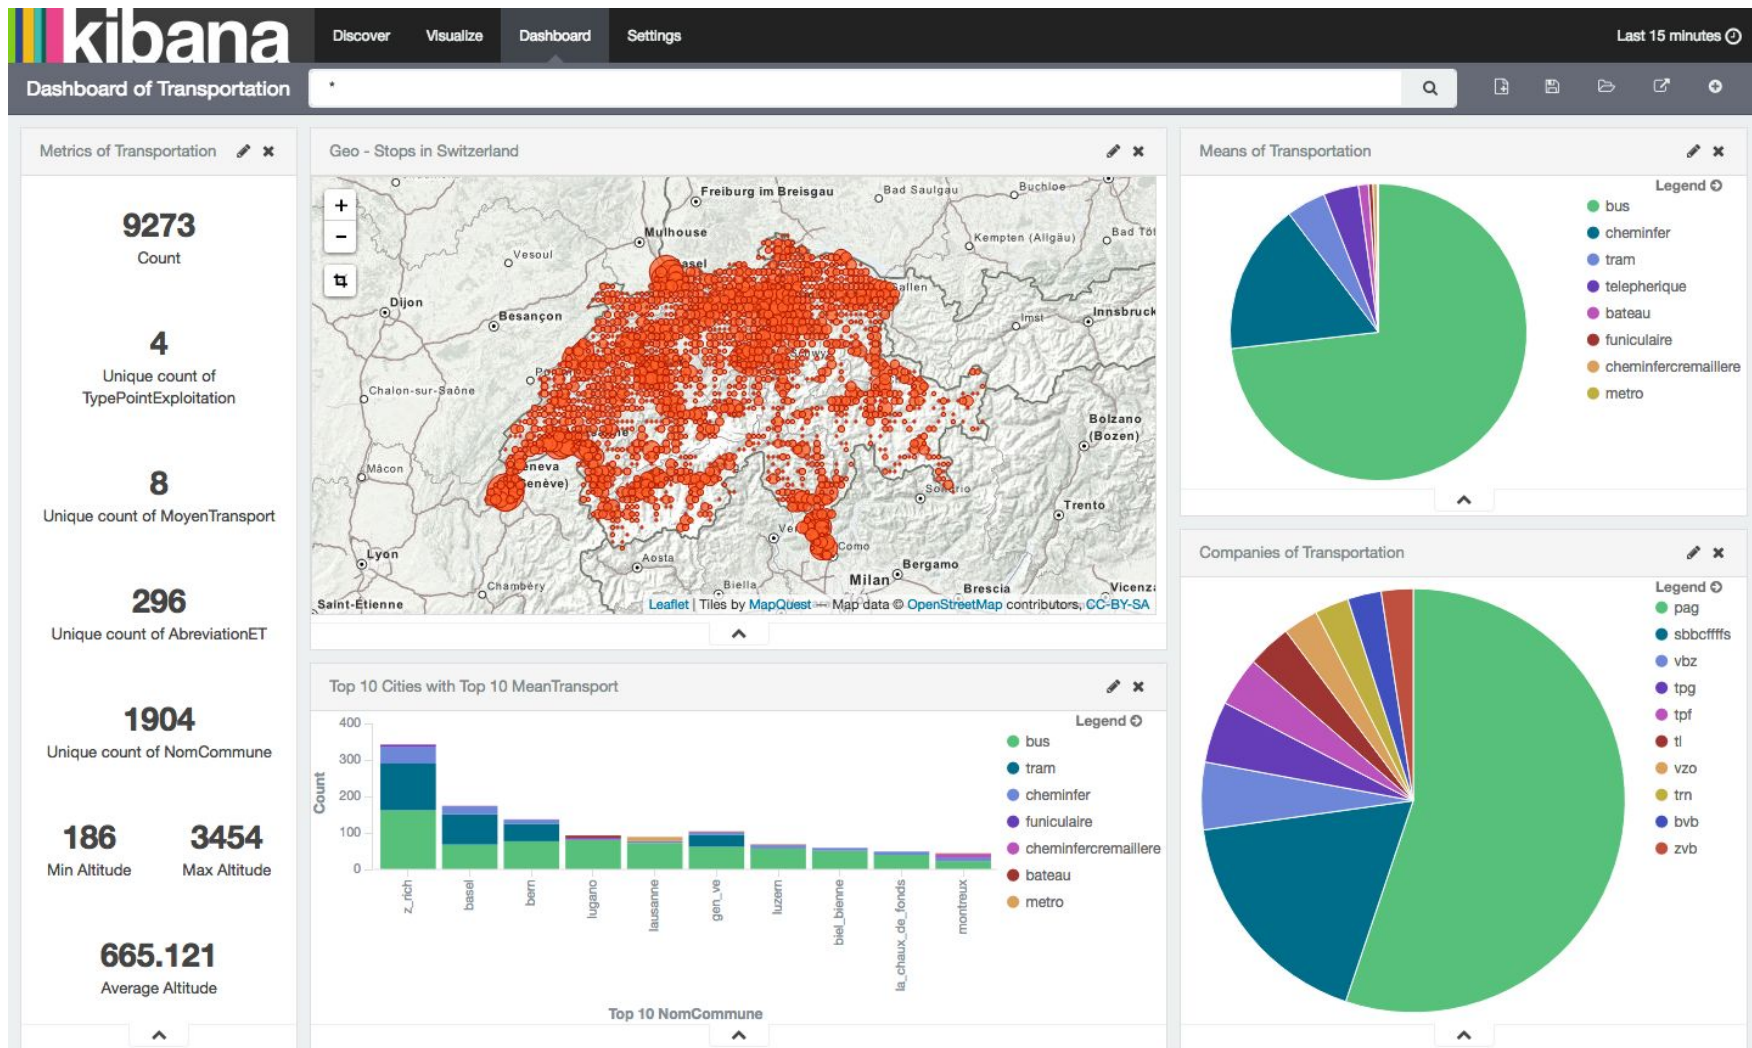
* <http://www.swisstopo.admin.ch/internet/swisstopo/fr/home/products/software/products/skripts.html>



Demo 2: Aggregation & Visualisation of Swiss Public Transportation Stops

- dashboard, basics metrics: number of unique ... by “field”
- pie charts, bar charts: count number of ... by “field”
- GPS visualisations
- cross-requests: Top 5 X by Y/ for Z
 - Top 10 Means by Top 10 Cities
 - Wrong Visualisation may hide data!
- Toggable fields & Make your own query !
 - Differences between `NomCommune:"Lausanne"` and `Nom:"Lausanne"` (analyzed fields)
 - `NomCommune:"Lausanne"` AND/OR NOT `MoyenTransport: Metro`
 - `_exists_:"var"` / `_missing_:"var"`
 - `+ term - term`

















Kibana - Dashboard - Swiss Public Transportation Stops Visualisations



Kibana - Global Metrics

[Discover](#)[Visualize](#)[Dashboard](#)[Settings](#)**transport****Metrics of Transportation**

metrics

- Metric** Count  
- Metric** Unique count of TypePointExploitation  
- Metric** Unique count of MoyenTransport  
- Metric** Unique count of AbreviationET  
- Metric** Unique count of NomCommune  
- Metric** Min Altitude  
- Metric** Max Altitude  
- Metric**  

Aggregation

Average 

Field

Altitude 

Advanced

+ Add Aggregation

view options

Apply

Discard

9273

Count

4

Unique count of TypePointExploitation

8

Unique count of MoyenTransport

296

Unique count of AbreviationET

*9273 Stops, of 4 Exploitation Type,
using 8 Means of Transport, exploited by 296 Companies,*

*in 1904 Municipalities, from an Altitude of 186 MASL (glitch, real is 197 in Locarno)
to 3454 MASL (Junfrauoch), with an Average Altitude of 665.1 MASL.*

1904

Unique count of NomCommune

186

Min Altitude

3454

Max Altitude

665.121

Average Altitude

Kibana - Pie Chart - Companies of Transportation

[Discover](#)[Visualize](#)[Dashboard](#)[Settings](#)

transport

metrics

▼ Slice Size

Aggregation

Count

◀ Advanced

buckets

▼ Split Slices

Aggregation

Terms

Field

⚠ Analyzed Field

AbbreviationET

Order

Size

Top

10

Order By

metric: Count

◀ Advanced

🔗 Add Sub Aggregation

view options ▶

Apply

Discard

Companies of Transportation

Legend

pag

sbbcofffs

vbz

tpg

tpf


tl

vzo

trn

bvb

zvb



Kibana - Pie Chart - Means of Transportation

[Discover](#)[Visualize](#)[Dashboard](#)[Settings](#)

transport

metrics

▼ Slice Size

Aggregation

Count

◀ Advanced

buckets

▼ Split Slices

Aggregation

Terms

Field

⚠ Analyzed Field

MoyenTransport

Order

Size

Top

10

Order By

metric: Count

◀ Advanced

⚡ Add Sub Aggregation

view options ▶

Apply

Discard

Means of Transportation

Legend ⓘ

● bus

● cheminfer

● tram


● telepherique

● bateau

● funiculaire

● cheminfercremailiere

● metro



Kibana - Tile Map (Geodata)



Discover

Visualize

Dashboard

Settings



transport

Geo - Stops in Switzerland

metrics

Value

Aggregation

Count

Advanced

buckets

Geo Coordinates

Aggregation

Geohash

Field

location

Precision

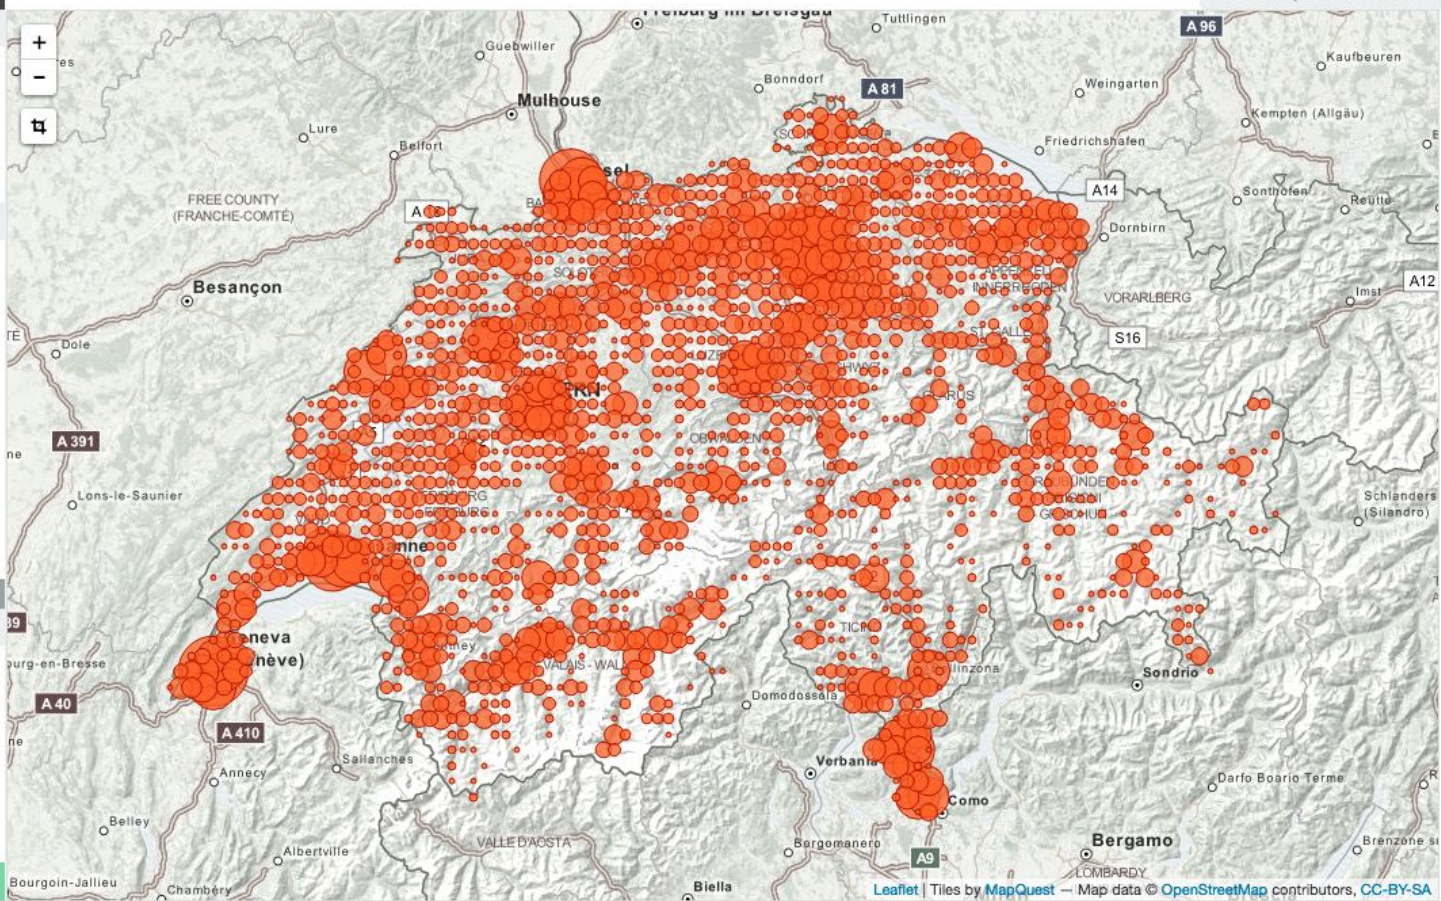
Advanced

Add Sub Aggregation

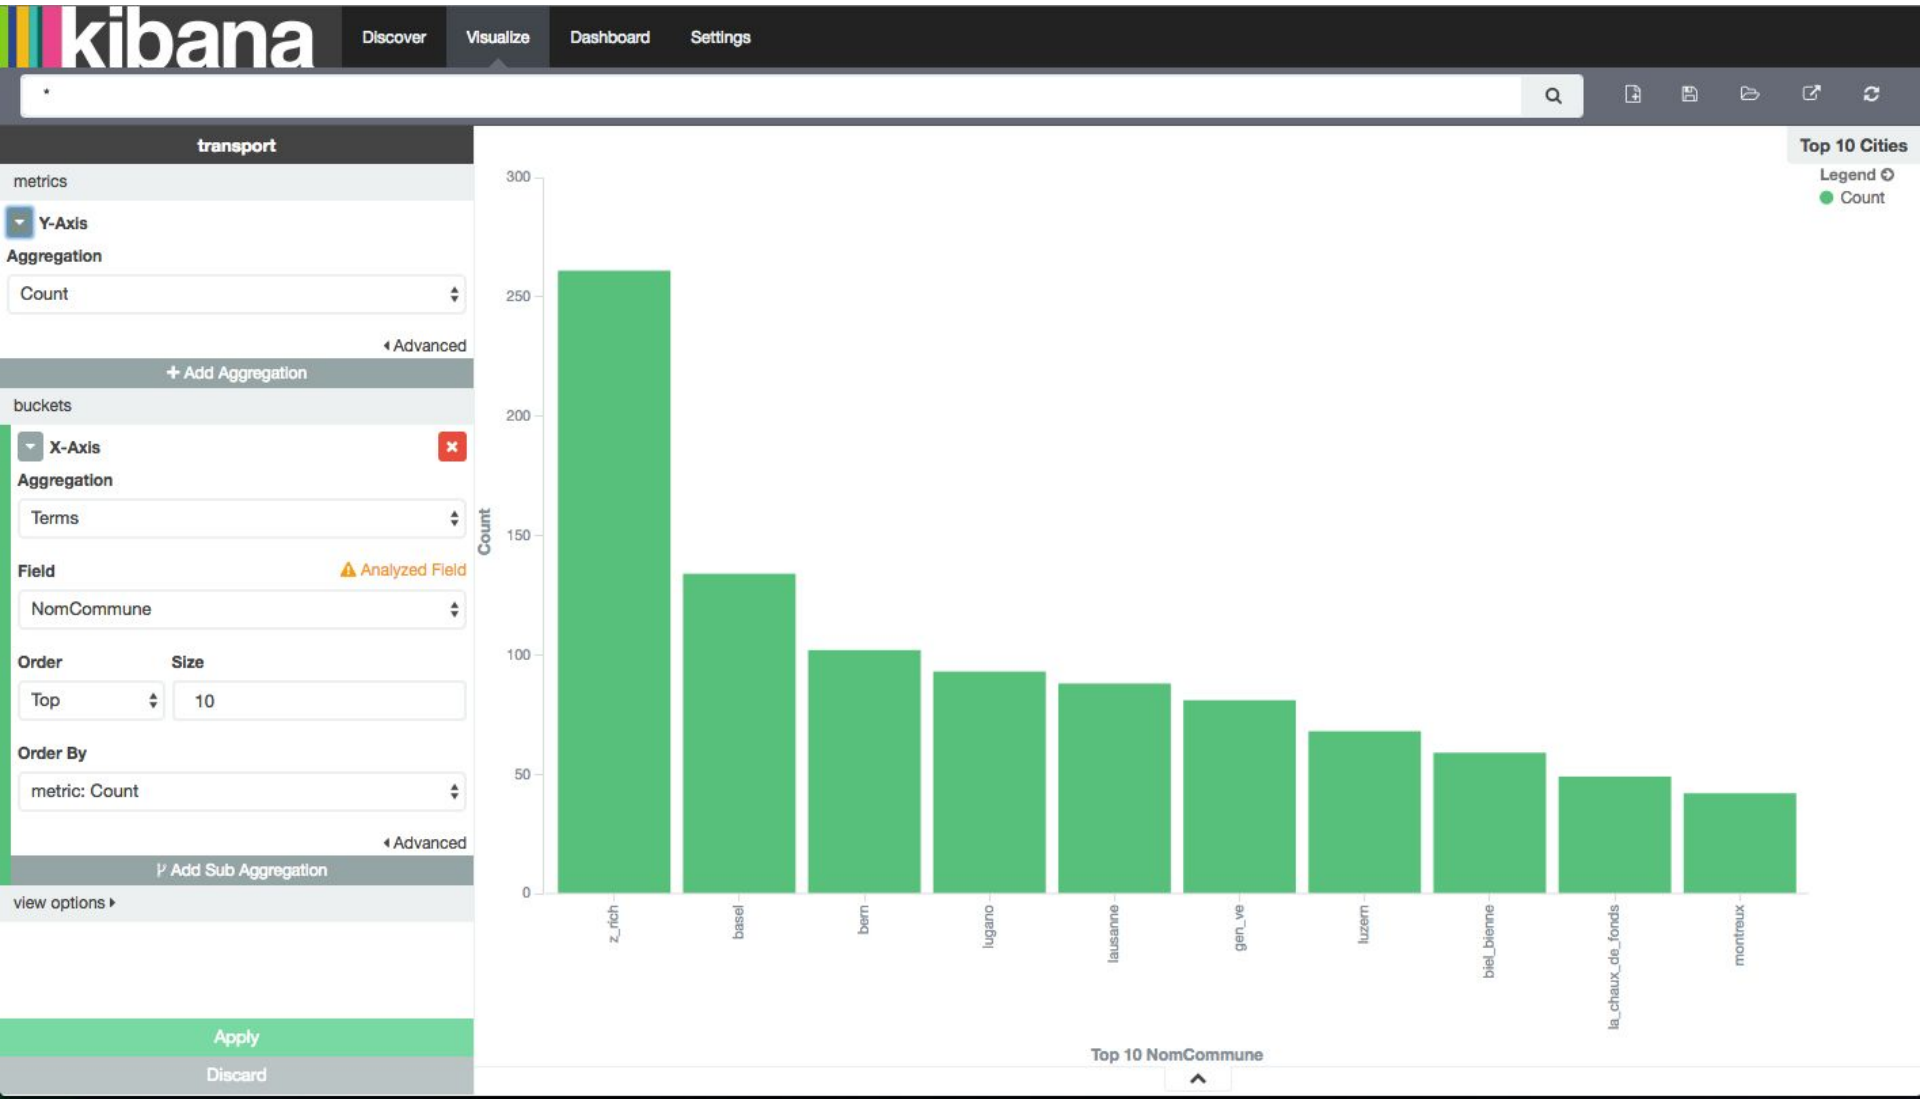
view options

Apply

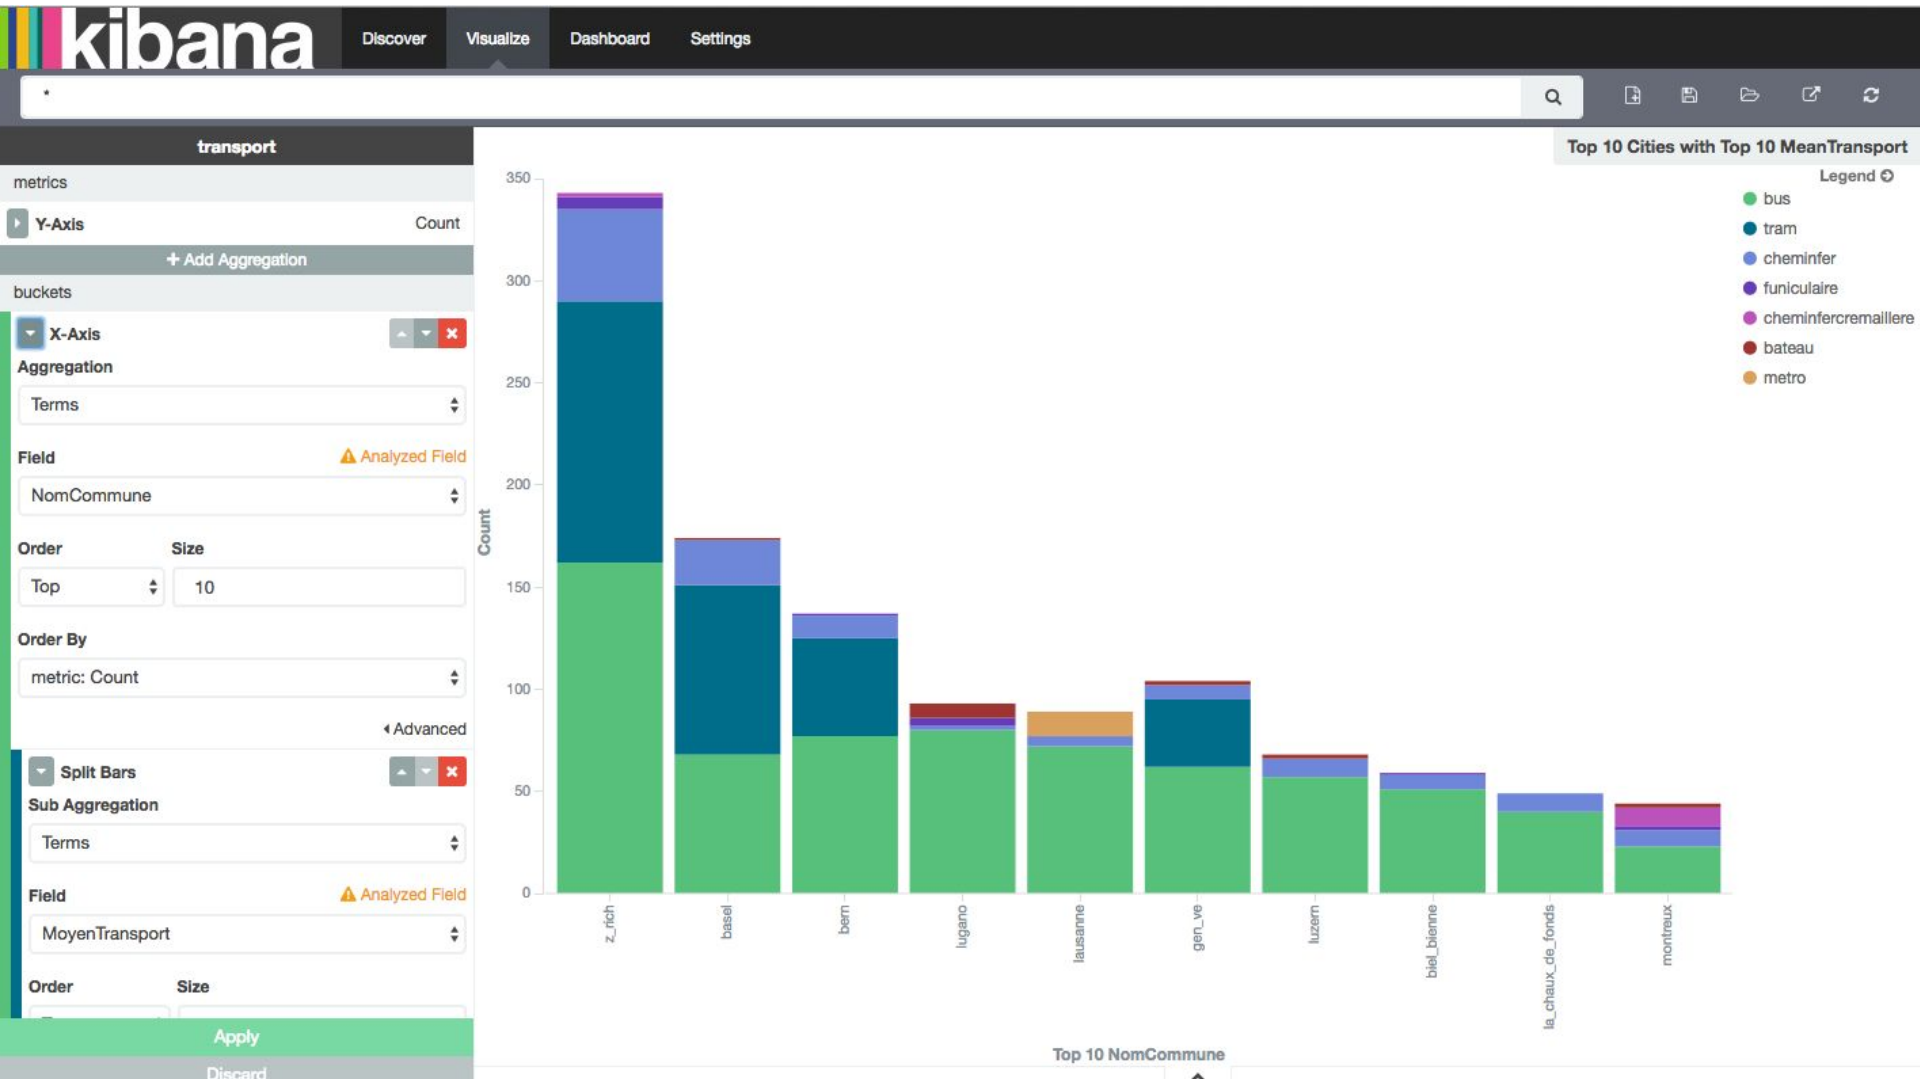
Discard



Kibana - Vertical Bar Chart - Top 10 Cities



Kibana - Vertical Bar Chart with Sub-Aggregation - Top 10 Means by Top 10 Cities



Kibana - Top 10 Means by Top 10 Companies VERSUS Top 10 Companies by Top 10 Means

[Discover](#)[Visualize](#)[Dashboard](#)[Settings](#)

transport

metrics

Y-Axis

Count

+ Add Aggregation

buckets

X-Axis

Aggregation

Terms

Field ⚠ Analyzed Field

AbreviationET

Order

Size

Top

10

Order By

metric: Count

⏮ Advanced

Split Bars

Sub Aggregation

Terms

Field ⚠ Analyzed Field

MoyenTransport

Order

Size


Apply

Discard

Top 10 Companies by Top 10 MeansTransport

Legend

- bus
- tram
- cheminifer
- funiculaire
- cheminifercremalliere
- metro



Kibana - Top 10 Means by Top 10 Companies VERSUS Top 10 Companies by Top 10 Means

[Discover](#)[Visualize](#)[Dashboard](#)[Settings](#)

transport

metrics

Y-Axis

Count

+ Add Aggregation

buckets

X-Axis

Aggregation

Terms

Field ⚠ Analyzed Field

MoyenTransport

Order

Size

Top

5

Order By

metric: Count

Advanced

Split Bars

Sub Aggregation

Terms

Field ⚠ Analyzed Field

AbreviationET

Order

Size

Apply

Discard

Top 5 MeansTransport by Top 10 Companies

Legend

pag

tpg

vbz

tpf

tl

vzo

zvb

trn

sbbcfffs

bos

fr

rhb

vr

bis

mgb

ab

tpc

asm

zb

bvb

blt

svb

rbs

vbg

aagl


fb

dkb

bfl

bdgag

hah



Kibana - Togglable Fields (OFF)

[Discover](#)[Visualize](#)[Dashboard](#)[Settings](#)

MoyenTransport: "tram" Actions ▶

transport

metrics

▶ Slice Size Count

buckets

▼ Split Slices ✕

Aggregation

Terms

Field ⚠ Analyzed Field

NomCommune

Order Size

Top 10

Order By

metric: Count

⏪ Advanced

⚡ Add Sub Aggregation

view options ▶


Apply

Discard

Top 10 Cities by Top 10 Companies (Tram toggle)

Legend ⓘ

- z_rich
- basel
- bern
- lugano
- lausanne
- gen_ve
- luzern
- biel_bienne
- la_chaux_de_fonds
- montreux



Kibana - Toggable Fields (ON)

[Discover](#)[Visualize](#)[Dashboard](#)[Settings](#)

MoyenTransport: "tram" [Actions](#)

transport

metrics

Slice Size Count

buckets

Split Slices ✕

Aggregation

Terms

Field ⚠ Analyzed Field

NomCommune

Order

Size

Top 10

Order By

metric: Count

⏪ Advanced

[Add Sub Aggregation](#)

[view options](#)


Apply

Discard

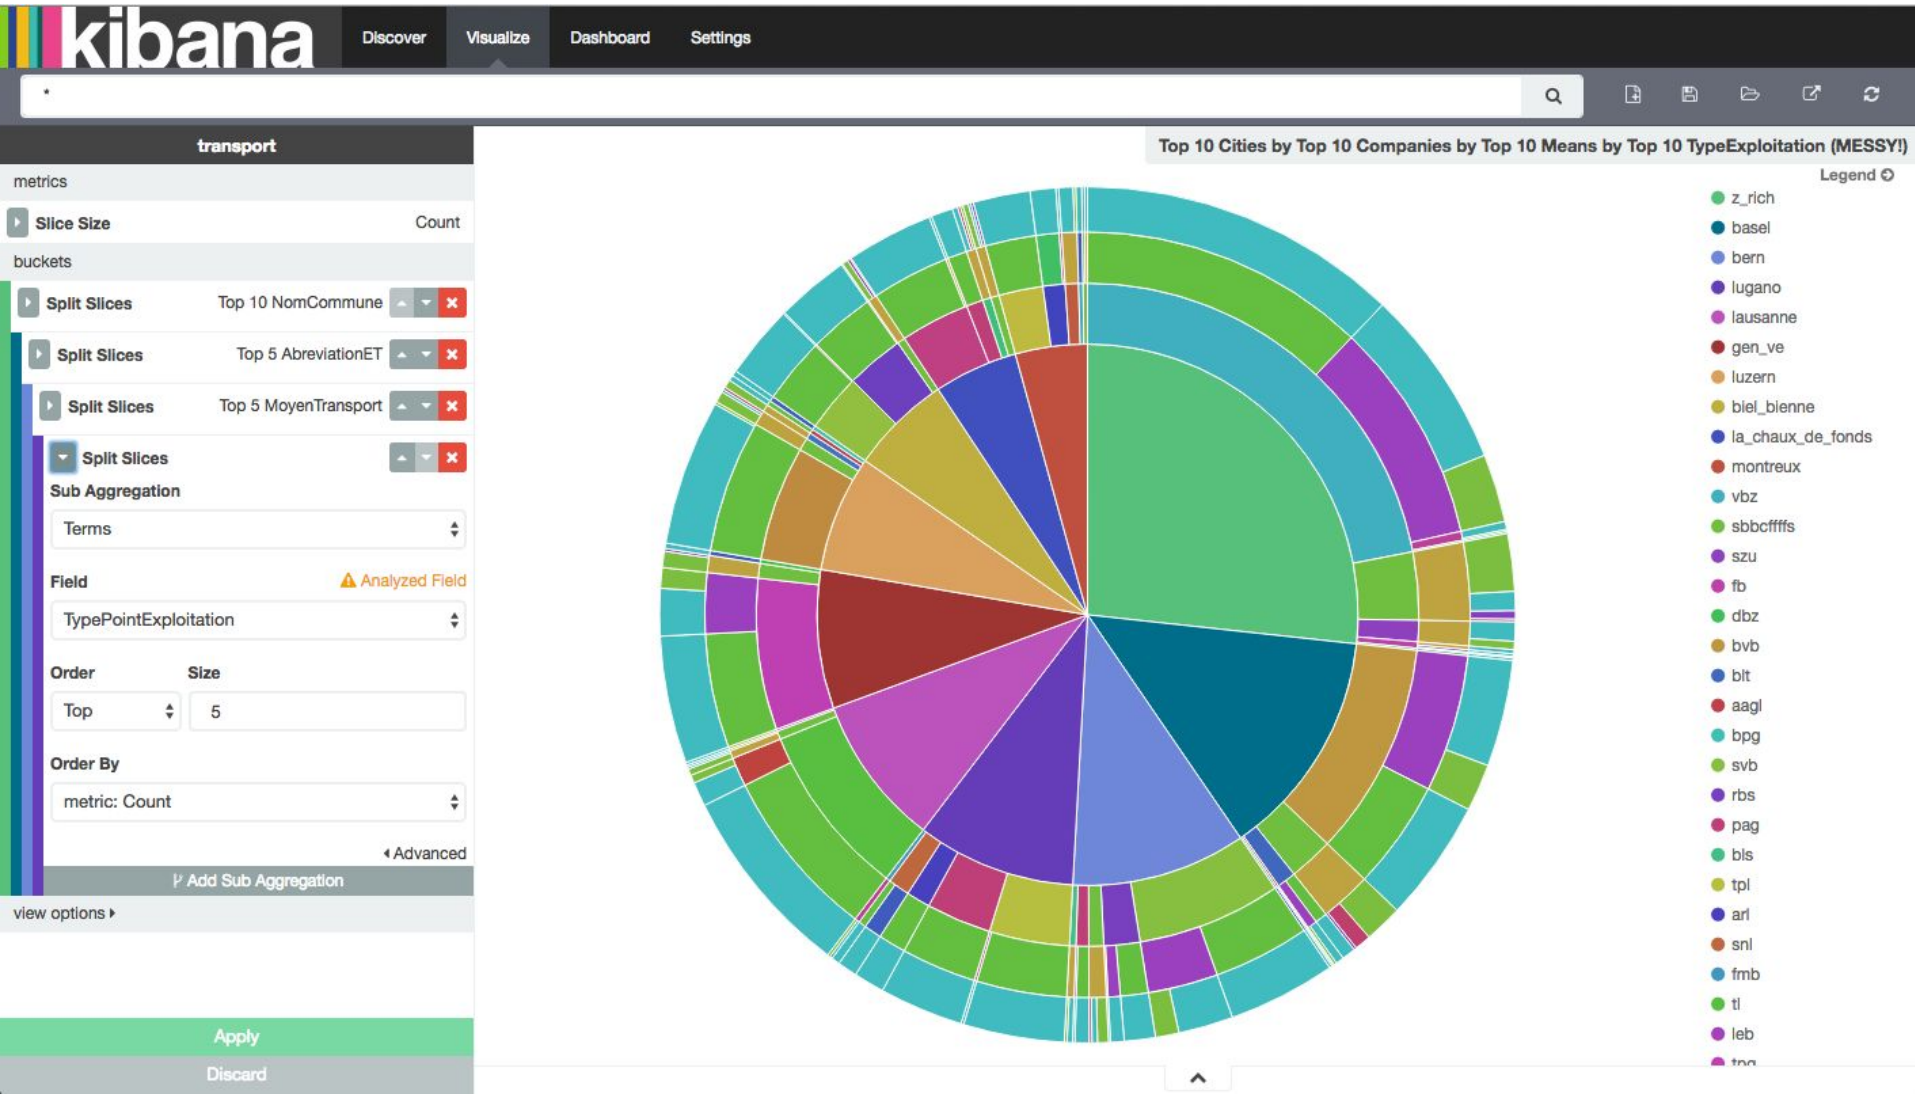
Top 10 Cities by Top 10 Companies (Tram toggle)

Legend

- z_rich
- basel
- bern
- gen_ve
- lancy
- relnach_bl_
- m_nchenstein
- meyrin
- riehen
- allschwil



Kibana - Keep it simple - Don't be messy !





Elasticsearch as an AWS Amazon Service ?

When building ELK for your business...

- option 1: deploy ELK in IT: Elastic, Logstash, Kibana = full stack solution
- option 2: use an ELK-as-a-service like Amazon or competitors



Elasticsearch as an AWS Amazon Service ?

Advantages ?

- All logs accessible in one place
- Nothing to install – use any log shipper you want
- Hassle-free managed ELK you don't need to maintain and scale
- easy to connect my existing services (IT or Amazon) to Amazon ELS
- Use built-in Kibana or use your own Kibana or Grafana with Logsene
- price: pay-as-you-go and per-service (no IT & infrastructure, maintenance...)
- price are low if the logs are from local Amazon source, transfer to or from the internet costs: see next slide.



Elasticsearch as an AWS Amazon Service ? Prices ?

- Machine: the cheapest is t2.micro.elasticsearch, 0.021\$/h, 15.12\$/month.
- storage: from \$0.079/GB/month (79 \$/TB/month) on magnetic disk (HDD)
- transfer of large amounts to or from the internet costs. All the log data have to go IN first: it's free from the internet, but might cost from certain services (up to \$ 0.01/GB, 10 \$/TB). Transferring data OUT to the internet is expensive (up to \$0.090 per GB, that's 90\$/TB).
- do all other computation also on an Amazon service (with free or cheap transfer rate). Usually Kibana will only transfer small amounts of data for what it needs to render the graphics or the search query.

<https://aws.amazon.com/elasticsearch-service/pricing/>



Elasticsearch as an AWS Amazon Service ?

Drawbacks (in favor of ELK Full-stack in IT dept.)

- version of Elasticsearch: not choosable, and Elasticsearch 1.5.x and other versions have critical bugs
- less flexible with library support, the technology moves fast!
- limited choices in terms of VM characteristics: (type, disk size, RAM)
- requires cloud usage (country and legal issues, technical issues with traffic/bandwidth)



Competitors of ELS on AWS ?

- **Logsene** by SemaText is an alternative service that integrates with SPM Performance Monitoring. Available in the Cloud and On Premise. Correlate logs with performance metrics via SPM. UI view of logs (kibana-like).
- **Whoosh** is developed in Python (ELS in Java). Not as powerful but useful if you code in python.
- **Apache Solr** needs a schema (not ELS), has bigger community. Better performances.

sematext



Solr 

<https://sematext.com/logsene/> <https://pypi.python.org/pypi/Whoosh/> <https://lucene.apache.org/solr/>

To go further

Amazon official documentation

<https://aws.amazon.com/elasticsearch-service>

<http://aws.amazon.com/en/documentation/elasticsearch-service/>

Elastic.co official documentation

<https://www.elastic.co/products/elasticsearch>

<https://www.elastic.co/guide/en/elasticsearch/reference/current/index.html>

<http://aws.amazon.com/en/documentation/elasticsearch-service/>

Other sources of information

<https://wikitech.wikimedia.org/wiki/Logstash>

<https://cloudacademy.com/amazon-web-services/advanced-techniques-for-aws-monitoring-metrics-and-logging-course/>

<https://github.com/elastic/elasticsearch>

<http://www.elasticsearchtutorial.com/elasticsearch-in-5-minutes.html>

Thanks! Any questions ?

Elasticsearch, Logstash & Kibana

Search & Analyze log files in clusters & Real Time

Berney Leonard, Bron Sacha, Minder Valentin, Salathe Fabien
HEIG-VD - CLD - 2016



elastic



logstash

